

1 Introduction

This is a specification of a simple scheduler and assembler. The system contains a set of registers and a block of memory. Processes can be created, with each containing a sequence of instructions that are executed on the system. The instruction format is a simplified format of the Intel x86 architecture. Processes are scheduled based on the credit system that is found in the Linux 2.0 kernel.

2 Stack

This specification was written as a test spec for the CZT project. As a result, there are parts that may appear to be specified in a strange way - this is to test out the tools on a large set of Z.

section *Stack* **parents** *standard_toolkit*

A generic stack.

$Stack[X] == [stack : seq X]$
 $InitStack[X] == [Stack[X] \mid stack = \emptyset]$

$PushStack[X]$
$\Delta Stack[X]$
$x? : X$
$stack' = stack \frown \langle x? \rangle$

$PopStack[X]$
$\Delta Stack[X]$
$x! : X$
$stack' \frown \langle x! \rangle = stack$

3 Definitions

section *Definitions* **parents** *standard_toolkit*

Firstly, we define some basic types and functions that are used throughout the specification.

Declarations	This section	Globally
unboxed items	3	3
axiomatic definitions	0	0
generic axiomatic defs.	0	0
schemas	0	0
generic schemas	4	4
Total	7	7

Table 1: Summary of Z declarations for Section 2.

singleton is the set of all sets whose size is less than or equal to 1. This is included only to have a generic axiom definition.

relation(*singleton* _)

[X]
<i>singleton</i> _ : $\mathbb{P}(\mathbb{P} X)$
$\forall s : \mathbb{P} X \bullet \text{singleton } s \Leftrightarrow \# s \leq 1$

The basic type of this system is a word, which specifically, is an unsigned octet. An unsigned word is used so references to memory etc a 1-relative.

WORD == 0 .. 255

Then, we define the size of the memory block, and give it a value for animation purposes.

<i>mem_size</i> : WORD
<i>mem_size</i> = 100

A *LABEL* is used to label instructions for *jump* instructions etc, although 'jump' hasn't been specified yet.

[*LABEL*]

Now we define the different instructions, as well as their operands. A *CONSTANT* is used both as a constant value, as well as a memory reference for load and store instructions.

$INST_NAME ::=$
 $add \mid sub \mid divide \mid mult \mid push \mid pop \mid load \mid store \mid loadConst \mid print$
 $OPERAND ::= AX \mid BX \mid CX \mid DX \mid constant\langle\langle WORD \rangle\rangle$
 $REGISTER == \{AX, BX, CX, DX\}$
 $CONSTANT == OPERAND \setminus REGISTER$

An instruction is specified as a instruction name, a sequence of operands, and optionally, a label.

<i>Instruction</i> $label : \mathbb{P} LABEL$ $name : INST_NAME$ $params : seq OPERAND$ <i>singleton label</i>

Declarations	This section	Globally
unboxed items	9	12
axiomatic definitions	1	1
generic axiomatic defs.	1	1
schemas	3	3
generic schemas	0	4
Total	14	21

Table 2: Summary of Z declarations for Section 3.

4 System

section *System parents Definitions, Stack*

The system consists of a set of registers, and a block of memory. There is also a buffer for displaying output.

$REGISTERS == REGISTER \rightarrow OPERAND$
 $MEMORY == 1 .. mem_size \mapsto WORD$

System

registers : *REGISTERS*
memory : *MEMORY*
output : seq *WORD*

Initially, all registers and memory hold the minimum *WORD* value. The output buffer is empty.

InitSystem

System

registers = {*r* : *REGISTER* • *r* ↦ *constant*(*min*(*WORD*))}
memory = {*m* : 1 .. *mem_size* • *m* ↦ *min*(*WORD*)}
output = ⟨⟩

The system can have arithmetic and memory instructions.

Arith_Inst == [*Instruction* | # *params* = 2 ∧ *params*(1) ∈ *REGISTER*]
Add_Inst == [*Arith_Inst* | *name* = *add*]
Sub_Inst == [*Arith_Inst* | *name* = *sub*]
Mult_Inst == [*Arith_Inst* | *name* = *mult*]
Div_Inst == [*Arith_Inst* | *name* = *divide*]

Memory_Inst == [*Instruction* | # *params* = 2 ∧ *params*(1) ∈ *REGISTER*
 ∧ *params*(2) ∈ *CONSTANT*]
Load_Inst == [*Memory_Inst* | *name* = *load*]
LoadConst_Inst == [*Memory_Inst* | *name* = *loadConst*]
Store_Inst == [*Memory_Inst* | *name* = *store*]

A print instruction prints the value of a register.

Print_Inst == [*Instruction* | # *params* = 1]

val maps constants to their value, and *dereference* dereferences the value of a register, transitively if required.

val : *CONSTANT* → *WORD*
dereference : *OPERAND* × *REGISTERS* → *WORD*

∀ *c* : *CONSTANT* •
 (∃ *n* : *WORD* • *c* = *constant*(*n*) ∧ *val*(*c*) = *n*)
 ∀ *a* : *OPERAND*; *r* : *REGISTERS* •
dereference(*a*, *r*) =
 if *a* ∈ *REGISTER* then *dereference*(*r*(*a*), *r*) else *val*(*a*)

The specification of the arithmetic instructions.

<i>Add</i>
Δ_{System} <i>Add_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 + o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Sub</i>
Δ_{System} <i>Sub_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 - o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Mult</i>
Δ_{System} <i>Mult_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 * o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Div</i>
Δ_{System} <i>Div_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 \text{ div } o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

The **load** operation loads a constant from memory. The second parameter is an index to the memory location from which the constant is loaded.

<i>Load</i>
$\Delta System$
<i>Load_Inst</i>
$\begin{aligned} &\exists o_2 == \text{val}(\text{params}(2)) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \\ &\quad \quad \{\text{params}(1) \mapsto \text{constant}(\text{memory}(o_2))\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

loadConst loads a constant into a register. The second parameter the constant to be loaded.

<i>Load_Const</i>
$\Delta System$
<i>LoadConst_Inst</i>
$\begin{aligned} &\exists o_2 == \text{val}(\text{params}(2)) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

Store the value of a register in memory.

<i>Store</i>
$\Delta System$
<i>Store_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{val}(\text{params}(2)) \bullet \\ &\quad \quad \text{memory}' = \text{memory} \oplus \{o_2 \mapsto o_1\} \\ &\text{registers}' = \text{registers} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Print</i>
$\Xi System$
<i>Print_Inst</i>
$\begin{aligned} &\text{output}' = \text{output} \frown \langle \text{dereference}(\text{params}(1), \text{registers}) \rangle \\ &\text{registers}' = \text{registers} \\ &\text{memory}' = \text{memory} \end{aligned}$

$$\begin{aligned}
Stack_Inst &== [Instruction \mid \# params = 1] \\
Push_Inst &== [Stack_Inst \mid name = push] \\
Pop_Inst &== [Stack_Inst \mid name = pop]
\end{aligned}$$

The specification of the stack instructions on the system.

$ \begin{aligned} &Push0 \\ &\Xi System \\ &PushStack[WORD] \\ &Push_Inst \end{aligned} $
$x? = dereference(params(1), registers)$

$ \begin{aligned} &Pop0 \\ &\Delta System \\ &PopStack[WORD] \\ &Pop_Inst \end{aligned} $
$ \begin{aligned} registers' &= registers \oplus \{params(1) \mapsto constant(x!)\} \\ memory' &= memory \\ output' &= output \end{aligned} $

$$\begin{aligned}
Push &== Push0 \upharpoonright [System; Stack[WORD]] \\
Pop &== Pop0 \upharpoonright [System; Stack[WORD]]
\end{aligned}$$

This executes an instruction on the on the system. $inst?$ is the instruction to execute, and $base?$ is the base memory value of the executing process. If the instruction is a **load** or **store** instruction, the memory reference must offset using the base value.

<i>exec_inst</i>
$\Delta System$
<i>inst?</i> : <i>Instruction</i>
<i>base?</i> : 1 .. <i>mem_size</i>
$\exists label : \mathbb{P} LABEL; name : INST_NAME; params : seq OPERAND \mid$ $label = inst?.label \wedge name = inst?.name \wedge$ $params = inst?.params \bullet$ $Add \vee Sub \vee Mult \vee Div \vee$ $Print \vee Load_Const \vee$ $name \in \{load, store\} \Rightarrow (\exists p : seq OPERAND \mid$ $p = \langle params(1),$ $constant(val(params(2)) + base?) \rangle \bullet$ $Load[p/params] \vee Store[p/params])$

Declarations	This section	Globally
unboxed items	19	31
axiomatic definitions	6	7
generic axiomatic defs.	0	1
schemas	70	73
generic schemas	0	4
Total	95	116

Table 3: Summary of Z declarations for Section 4.

5 Scheduler

section *Scheduler* parents *System*

This part of the specification is the scheduler.

Here, we declare the set of process IDs, the priority values, and the default number of credits a process receives when it is created.

Pid == \mathbb{N}
Priority == - 19 .. 19
Default_Credits == 10

The possible status that a process can hold.

Status ::= *pWaiting* | *pReady* | *pRunning*

A process consists of a process ID, a status, a number of credits, and a priority. Each process has a sequence of instructions to be executed on the assembler, with a pointer to the current instruction. The memory that a process can occupy is between a base and limit value. Instructions must only access memory with a value less than the limit, but they know nothing about the base value - this is added onto the memory index provided by the instruction when an instruction is executed. Each process also contains a stack and values for all registers, which are used to store values when the process is suspended.

Processes

$$\begin{array}{l}
pids : \mathbb{P} \text{ Pid} \\
status : \text{Pid} \leftrightarrow \text{Status} \\
credits : \text{Pid} \leftrightarrow \mathbb{N} \\
priority : \text{Pid} \leftrightarrow \text{Priority} \\
instructions : \text{Pid} \leftrightarrow (\text{seq } \text{Instruction}) \\
inst_pointer : \text{Pid} \leftrightarrow \mathbb{N}_1 \\
base, limit : \text{Pid} \leftrightarrow \text{WORD} \\
pregisters : \text{Pid} \leftrightarrow \text{REGISTERS} \\
pstack : \text{Pid} \leftrightarrow \text{Stack}[\text{WORD}] \\
\\
pids = \text{dom}(status) = \text{dom}(credits) = \text{dom}(priority) = \\
\quad \text{dom}(instructions) = \text{dom}(inst_pointer) = \text{dom}(base) = \\
\quad \text{dom}(limit) = \text{dom}(pstack) \\
\forall pid : pids \bullet inst_pointer(pid) \leq \#(instructions(pid)) \\
\forall pid : pids \bullet base(pid) + limit(pid) \leq mem_size
\end{array}$$

The *sort* function takes the credits and priorities of all processes, and returns a sequence of process IDS sorted firstly by their credits (the more credits a process has, the higher preference they get), and if the credits are equal, then their priority. If the priority is equal, then the order is non-deterministic.

$$\begin{array}{l}
sort : (\text{Pid} \leftrightarrow \mathbb{N}) \times (\text{Pid} \leftrightarrow \text{Priority}) \leftrightarrow \text{iseq } \text{Pid} \\
\\
sort = (\lambda credits : (\text{Pid} \leftrightarrow \mathbb{N}); priority : (\text{Pid} \leftrightarrow \text{Priority}) \mid \\
\quad \text{dom}(credits) = \text{dom}(priority) \bullet \\
\quad (\mu s : \text{iseq } \text{Pid} \mid \text{ran}(s) = \text{dom}(credits) \wedge \\
\quad \quad (\forall i : 1 \dots \# s - 1 \bullet \\
\quad \quad \quad credits(s(i)) > credits(s(i+1)) \vee \\
\quad \quad \quad (credits(s(i)) = credits(s(i+1))) \wedge \\
\quad \quad \quad priority(s(i)) > priority(s(i+1)))) \bullet s))
\end{array}$$

To interrupt a process during execution, the kernel must be in *kernel* mode.

Mode ::= user | kernel

For the scheduler, we track which mode the operating system is in, as well as declaring three “secondary” variables, *waiting*, *running*, and *ready*, to keep the sets of waiting running, and ready variables respectively. In fact, *ready* is a sequence, and is ordered based on the credits that each process has. A process with more credits will have a higher priority. This is fair scheduling, because at each timer interrupt (the *tick* operation specified below), the current process losses one credit, therefore, process spending a lot of time executing will eventually have a low priority.

<i>Scheduler</i>
<i>Processes</i>
<i>System</i>
<i>Stack</i> [<i>WORD</i>]
<i>mode</i> : <i>Mode</i>
<i>waiting, running</i> : $\mathbb{P} \text{ Pid}$
<i>ready</i> : <i>iseq Pid</i>
$\# \text{ running} \leq 1$
$\text{waiting} \cap \text{running} \cap \text{ranready} = \emptyset$
$\text{waiting} \cup \text{running} \cup \text{ranready} = \text{pids}$
$\text{waiting} = \{p : \text{pids} \mid (\text{status} \sim)(p\text{Waiting}) = p\}$
$\text{running} = \{p : \text{pids} \mid (\text{status} \sim)(p\text{Running}) = p\}$
$\text{ready} = \text{sort}((\text{waiting} \cup \text{running}) \triangleleft \text{credits},$
$\quad (\text{waiting} \cup \text{running}) \triangleleft \text{priority})$
$\forall r : \text{ran}(\text{ready}) \bullet \text{status}(r) = p\text{Ready}$
$\forall r : \text{running} \bullet \text{credits}(r) > 0$

This uses semicolons as conjunctions for predicates, which conforms to the grammar in the ISO standard, but according to the list of differences between ZRM and ISO Z on Ian Toyn’s website, semicolons can no longer be used to conjoin predicates.

<i>InitScheduler</i>
<i>Scheduler</i>
<i>InitStack</i> [<i>WORD</i>]
<i>InitSystem</i>
$\text{pids} = \emptyset; \text{status} = \emptyset; \text{priority} = \emptyset$
$\text{credits} = \emptyset; \text{instructions} = \emptyset; \text{inst_pointer} = \emptyset$
$\text{waiting} = \emptyset; \text{running} = \emptyset; \text{ready} = \langle \rangle$
$\text{base} = \{\}; \text{limit} = \{\}; \text{pregisters} = \{\}$
$\text{mode} = \text{user}$

newProcess creates a new process with a unique process ID and a specified priority, and places this new process on the ready queue.

<i>create_new_process</i>	<hr/> $\Delta Scheduler$ $\Xi System$ $priority? : Priority$ $instructions? : seq\ Instruction$ $base?, limit? : WORD$ $pid! : Pid$ <hr/> $pid! \notin pids$ $status' = status \cup \{pid! \mapsto pReady\}$ $credits' = credits \cup \{pid! \mapsto Default_Credits\}$ $priority' = priority \cup \{pid! \mapsto priority?\}$ $instructions' = instructions \cup \{pid! \mapsto instructions?\}$ $inst_pointer' = inst_pointer \cup \{pid! \mapsto 1\}$ $base' = base \cup \{pid! \mapsto base?\}$ $limit' = limit \cup \{pid! \mapsto limit?\}$ $pregisters' =$ $\quad pregisters \cup \{pid! \mapsto \{r : REGISTER \bullet$ $\quad \quad r \mapsto constant(min(WORD))\}\}$ $pstack' = pstack \cup \{pid! \mapsto (\langle stack == \rangle)\}$ $pids' = pids \cup \{pid!\}$ <hr/>
---------------------------	--

We define a schema that contains only the variables that do not change when a reschedule occurs.

$$RescheduleChange == \\ Scheduler \setminus (status, running, ready, waiting, credits)$$

A reschedule occurs when all ready processes have no credits. Every process, not just the ready processes, have their credits re-calculated using the formula $credits = credits/2 + priority$. This implies that the ready process with the highest priority will be the next process executed.

<i>reschedule</i>	<hr/> $\Delta Scheduler$ $\Xi RescheduleChange$ <hr/> $ready \neq \emptyset$ $\forall r : ran(ready) \bullet credits(r) = 0 \Rightarrow$ $\quad credits' = \{p : pids \bullet p \mapsto (credits(p) \div 2) + priority(p)\} \wedge$ $\quad status' = status$ $\neg (\forall r : ran(ready) \bullet credits(r) = 0) \Rightarrow$ $\quad status' = status \oplus \{head(ready) \mapsto pRunning\} \wedge$ $\quad credits' = credits$ <hr/>
-------------------	---

We declare a new schema that contains only the state variables that do not change when a status change occurs.

$$\text{StatusChange} == \text{Scheduler} \setminus (\text{status}, \text{running}, \text{waiting}, \text{ready}, \text{registers}, \text{pregisters}, \text{pstack})$$

Interrupts the currently executing process if the new process is of a higher priority then the current process and the kernel is in *kernel* mode.

$$\begin{array}{l} \text{interrupt} \text{ ---} \\ \Delta \text{Scheduler} \\ \exists \text{StatusChange} \\ \text{create_new_process} \\ \hline \text{mode} = \text{kernel} \\ \text{running} = \emptyset \vee (\exists p : \text{running} \bullet \text{priority?} \geq \text{priority}(p)) \\ \exists r : \text{running} \bullet \\ \quad \text{status}' = \text{status} \oplus \{pid! \mapsto p\text{Running}, r \mapsto p\text{Ready}\} \wedge \\ \quad \text{pregisters}' = \text{pregisters} \oplus \{r \mapsto \text{registers}\} \wedge \\ \quad \theta \text{Stack}' = \text{pstack}(r) \\ \quad \text{registers}' = \text{pregisters}(pid!) \end{array}$$

Remove the currently running process and put it back in the ready queue.

$$\begin{array}{l} \text{remove_running_process} \text{ ---} \\ \Delta \text{Scheduler} \\ \exists \text{StatusChange} \\ \hline \exists pid == (\mu r : \text{running}) \bullet \\ \quad \text{status}' = \text{status} \oplus \{pid \mapsto p\text{Ready}\} \wedge \\ \quad \text{pregisters}' = \text{pregisters} \oplus \{pid \mapsto \text{registers}\} \wedge \\ \quad \text{pstack}' = \text{pstack} \oplus \{pid \mapsto \theta \text{Stack}'\} \end{array}$$

A process becomes blocked if it is waiting on a resource such as an IO device, or waiting on another process

$$\text{block_process} == \text{remove_running_process} \circ \text{reschedule}$$

We declare a schema containing only the variables that change for an unblock.

$$\text{UnblockProcessChange} == \text{Scheduler} \setminus (\text{status}, \text{running}, \text{ready}, \text{waiting})$$

Unblocks a process that is blocked by another process.

unblock_process <hr/> $\Delta \text{Scheduler}$ $\Xi \text{UnblockProcessChange}$ $pid? : \text{Pid}$ <hr/> $pid? \in \text{pids}$ $\text{status}(pid?) = p \text{Waiting}$ $\text{running} = \emptyset \Leftrightarrow \text{status}' = \text{status} \oplus \{pid? \mapsto p \text{Running}\}$ $\text{running} \neq \emptyset \Leftrightarrow \text{status}' = \text{status} \oplus \{pid? \mapsto p \text{Ready}\}$ <hr/>
--

Remove a process from the system

remove_process <hr/> $\Delta \text{Scheduler}$ $\Xi \text{Stack}[\text{WORD}]$ ΞSystem $pid? : \text{Pid}$ <hr/> $pid? \in \text{pids}$ $\text{pids}' = \text{pids} \setminus \{pid?\}$ $\text{status}' = \{pid?\} \triangleleft \text{status}$ $\text{credits}' = \{pid?\} \triangleleft \text{credits}$ $\text{priority}' = \{pid?\} \triangleleft \text{priority}$ $\text{instructions}' = \{pid?\} \triangleleft \text{instructions}$ $\text{inst_pointer}' = \{pid?\} \triangleleft \text{inst_pointer}$ $\text{base}' = \{pid?\} \triangleleft \text{base}$ $\text{limit}' = \{pid?\} \triangleleft \text{limit}$ $\text{pregisters}' = \{pid?\} \triangleleft \text{pregisters}$ $\text{pstack}' = \{pid?\} \triangleleft \text{pstack}$ <hr/>

Update the details in the process table when each instruction is executed, as well as communicate the current instruction and the base value for the current process.

$$\text{ChangeInstPointer} == \text{Scheduler} \setminus (\text{inst_pointer})$$

<i>update_process_table</i>
$\Delta \text{Scheduler}$ $inst! : \text{Instruction}$ $base! : \text{WORD}$
$running \neq \emptyset$ $(\exists pid == (\mu r : running) \bullet$ $\quad inst! = head(instructions(pid)) \wedge$ $\quad base! = base(pid) \wedge$ $\quad (inst_pointer(pid) = \#(instructions(pid)) \Rightarrow$ $\quad \quad remove_process[pid/pid?]) \wedge$ $\quad inst_pointer(pid) < \#(instructions(pid)) \Rightarrow$ $\quad \quad inst_pointer' =$ $\quad \quad \quad inst_pointer \oplus \{pid \mapsto inst_pointer(pid) + 1\})$ $\quad \theta \text{ ChangeInstPointer} = \theta \text{ ChangeInstPointer}'$

$next == exec_inst \gg update_process_table \circ$
 $([\Delta \text{Scheduler} \mid running = \emptyset] \wedge reschedule) \vee$
 $([\Xi \text{Scheduler} \mid running \neq \emptyset])$

$idle0 == \neg \text{pre } next$

<i>idle</i>
$\Xi \text{Scheduler}$ $inst? : \text{Instruction}$ $base? : \text{WORD}$
$idle0$

$tick == next \vee idle$

$\vdash? (\forall n : \mathbb{N}_1 \bullet n > 0)$

$[X] \vdash? \forall x : \mathbb{P} X \bullet \# x \leq 1 \Leftrightarrow singleton\ x$

theorem PreconditionCheck
 $\forall \text{Scheduler} \bullet \text{pre } update_process_table$

Declarations	This section	Globally
unboxed items	21	52
axiomatic definitions	7	14
generic axiomatic defs.	0	1
schemas	110	183
generic schemas	0	4
Total	138	254

Table 4: Summary of Z declarations for Section 5.