

1 Introduction

This is a specification of a simple scheduler and assembler. The system contains a set of registers and a block of memory. Processes can be created, with each containing a sequence of instructions that are executed on the system. The instruction format is a simplified format of the Intel x86 architecture. Processes are scheduled based on the credit system that is found in the Linux 2.0 kernel.

2 Stack

This specification was written as a test spec for the CZT project. As a result, there are parts that may appear to be specified in a strange way - this is to test out the tools on a large set of Z.

section *Stack* **parents** *standard_toolkit*

A generic stack.

$Stack[X] == [stack : seq X]$
 $InitStack[X] == [Stack[X] \mid stack = \emptyset]$

| |
|-------------------------------------------------|
| $PushStack[X]$ $\Delta Stack[X]$ $x? : X$ |
| $stack' = stack \frown \langle x? \rangle$ |

| |
|------------------------------------------------|
| $PopStack[X]$ $\Delta Stack[X]$ $x! : X$ |
| $stack' \frown \langle x! \rangle = stack$ |

Ok, lets see the value of 3 unboxed items in Section 2!

3 Definitions

section *Definitions* **parents** *standard_toolkit*

| Declarations | This section | Globally |
|-------------------------|--------------|----------|
| Unboxed items | 3 | 3 |
| Axiomatic definitions | 0 | 0 |
| Generic axiomatic defs. | 0 | 0 |
| Schemas | 0 | 0 |
| Generic schemas | 2 | 2 |
| Total | 5 | 5 |

Table 1: Summary of Z declarations for Section 2.

Firstly, we define some basic types and functions that are used throughout the specification.

singleton is the set of all sets whose size is less than or equal to 1. This is included only to have a generic axiom definition.

relation(*singleton* _)

| | |
|------------------------------------------------------------------------------------|--|
| [X] | |
| <i>singleton</i> _ : $\mathbb{P}(\mathbb{P} X)$ | |
| $\forall s : \mathbb{P} X \bullet \text{singleton } s \Leftrightarrow \# s \leq 1$ | |

The basic type of this system is a word, which specifically, is an unsigned octet. An unsigned word is used so references to memory etc a 1-relative.

WORD == 0 .. 255

Then, we define the size of the memory block, and give it a value for animation purposes.

| |
|-------------------------------|
| <i>mem_size</i> : WORD |
| <i>mem_size</i> = 100 |

A *LABEL* is used to label instructions for *jump* instructions etc, although 'jump' hasn't been specified yet.

[*LABEL*]

Now we define the different instructions, as well as their operands. A *CONSTANT* is used both as a constant value, as well as a memory reference for load and store instructions.

$INST_NAME ::=$
 $add \mid sub \mid divide \mid mult \mid push \mid pop \mid load \mid store \mid loadConst \mid print$
 $OPERAND ::= AX \mid BX \mid CX \mid DX \mid constant\langle\langle WORD \rangle\rangle$
 $REGISTER == \{AX, BX, CX, DX\}$
 $CONSTANT == OPERAND \setminus REGISTER$

An instruction is specified as a instruction name, a sequence of operands, and optionally, a label.

| |
|---------------------------------------------------------------------------------------------------------------------------------|
| <i>Instruction</i> $label : \mathbb{P} LABEL$ $name : INST_NAME$ $params : seq OPERAND$ <i>singleton label</i> |
|---------------------------------------------------------------------------------------------------------------------------------|

| Declarations | This section | Globally |
|-------------------------|--------------|-----------|
| Unboxed items | 9 | 12 |
| Axiomatic definitions | 1 | 1 |
| Generic axiomatic defs. | 1 | 1 |
| Schemas | 1 | 1 |
| Generic schemas | 0 | 2 |
| Total | 12 | 17 |

Table 2: Summary of Z declarations for Section 3.

4 System

section *System parents Definitions, Stack*

The system consists of a set of registers, and a block of memory. There is also a buffer for displaying output.

$REGISTERS == REGISTER \rightarrow OPERAND$
 $MEMORY == 1 .. mem_size \leftrightarrow WORD$

System

registers : *REGISTERS*
memory : *MEMORY*
output : seq *WORD*

Initially, all registers and memory hold the minimum *WORD* value. The output buffer is empty.

InitSystem

System

registers = {*r* : *REGISTER* • *r* \mapsto *constant*(*min*(*WORD*))}
memory = {*m* : 1 .. *mem_size* • *m* \mapsto *min*(*WORD*)}
output = $\langle \rangle$

The system can have arithmetic and memory instructions.

Arith_Inst == [*Instruction* | # *params* = 2 \wedge *params*(1) \in *REGISTER*]
Add_Inst == [*Arith_Inst* | *name* = *add*]
Sub_Inst == [*Arith_Inst* | *name* = *sub*]
Mult_Inst == [*Arith_Inst* | *name* = *mult*]
Div_Inst == [*Arith_Inst* | *name* = *divide*]

Memory_Inst == [*Instruction* | # *params* = 2 \wedge *params*(1) \in *REGISTER*
 \wedge *params*(2) \in *CONSTANT*]
Load_Inst == [*Memory_Inst* | *name* = *load*]
LoadConst_Inst == [*Memory_Inst* | *name* = *loadConst*]
Store_Inst == [*Memory_Inst* | *name* = *store*]

A print instruction prints the value of a register.

Print_Inst == [*Instruction* | # *params* = 1]

val maps constants to their value, and *dereference* dereferences the value of a register, transitively if required.

val : *CONSTANT* \rightarrow *WORD*
dereference : *OPERAND* \times *REGISTERS* \rightarrow *WORD*

$\forall c : \text{CONSTANT} \bullet$
 $(\exists n : \text{WORD} \bullet c = \text{constant}(n) \wedge \text{val}(c) = n)$
 $\forall a : \text{OPERAND}; r : \text{REGISTERS} \bullet$
 $\text{dereference}(a, r) =$
if *a* \in *REGISTER* **then** *dereference*(*r*(*a*), *r*) **else** *val*(*a*)

The specification of the arithmetic instructions.

| |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Add</i> |
| Δ_{System} Add_Inst |
| $\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 + o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$ |

| |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Sub</i> |
| Δ_{System} Sub_Inst |
| $\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 - o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$ |

| |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Mult</i> |
| Δ_{System} Mult_Inst |
| $\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 * o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$ |

| |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Div</i> |
| Δ_{System} Div_Inst |
| $\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 \text{ div } o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$ |

The **load** operation loads a constant from memory. The second parameter is an index to the memory location from which the constant is loaded.

| |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Load</i> |
| $\Delta System$ |
| <i>Load_Inst</i> |
| $\begin{aligned} &\exists o_2 == \text{val}(\text{params}(2)) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \\ &\quad \quad \{\text{params}(1) \mapsto \text{constant}(\text{memory}(o_2))\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$ |

loadConst loads a constant into a register. The second parameter the constant to be loaded.

| |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Load_Const</i> |
| $\Delta System$ |
| <i>LoadConst_Inst</i> |
| $\begin{aligned} &\exists o_2 == \text{val}(\text{params}(2)) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$ |

Store the value of a register in memory.

| |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Store</i> |
| $\Delta System$ |
| <i>Store_Inst</i> |
| $\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{val}(\text{params}(2)) \bullet \\ &\quad \quad \text{memory}' = \text{memory} \oplus \{o_2 \mapsto o_1\} \\ &\text{registers}' = \text{registers} \\ &\text{output}' = \text{output} \end{aligned}$ |

| |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Print</i> |
| $\Xi System$ |
| <i>Print_Inst</i> |
| $\begin{aligned} &\text{output}' = \text{output} \frown \langle \text{dereference}(\text{params}(1), \text{registers}) \rangle \\ &\text{registers}' = \text{registers} \\ &\text{memory}' = \text{memory} \end{aligned}$ |

$$\begin{aligned}
Stack_Inst &== [Instruction \mid \# params = 1] \\
Push_Inst &== [Stack_Inst \mid name = push] \\
Pop_Inst &== [Stack_Inst \mid name = pop]
\end{aligned}$$

The specification of the stack instructions on the system.

| |
|--------------------------------------------------------------------------------------------|
| $ \begin{aligned} &Push0 \\ &\Xi System \\ &PushStack[WORD] \\ &Push_Inst \end{aligned} $ |
| $x? = dereference(params(1), registers)$ |

| |
|-----------------------------------------------------------------------------------------------------------------------------------------------|
| $ \begin{aligned} &Pop0 \\ &\Delta System \\ &PopStack[WORD] \\ &Pop_Inst \end{aligned} $ |
| $ \begin{aligned} registers' &= registers \oplus \{params(1) \mapsto constant(x!)\} \\ memory' &= memory \\ output' &= output \end{aligned} $ |

$$\begin{aligned}
Push &== Push0 \upharpoonright [System; Stack[WORD]] \\
Pop &== Pop0 \upharpoonright [System; Stack[WORD]]
\end{aligned}$$

This executes an instruction on the on the system. $inst?$ is the instruction to execute, and $base?$ is the base memory value of the executing process. If the instruction is a **load** or **store** instruction, the memory reference must offset using the base value.

| |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>exec_inst</i> |
| $\Delta System$ |
| <i>inst?</i> : <i>Instruction</i> |
| <i>base?</i> : 1 .. <i>mem_size</i> |
| $\exists label : \mathbb{P} LABEL; name : INST_NAME; params : seq OPERAND \mid$ $label = inst?.label \wedge name = inst?.name \wedge$ $params = inst?.params \bullet$ $Add \vee Sub \vee Mult \vee Div \vee$ $Print \vee Load_Const \vee$ $name \in \{load, store\} \Rightarrow (\exists p : seq OPERAND \mid$ $p = \langle params(1),$ $constant(val(params(2)) + base?) \rangle \bullet$ $Load[p/params] \vee Store[p/params])$ |

| Declarations | This section | Globally |
|-------------------------|--------------|-----------|
| Unboxed items | 19 | 31 |
| Axiomatic definitions | 1 | 2 |
| Generic axiomatic defs. | 0 | 1 |
| Schemas | 13 | 14 |
| Generic schemas | 0 | 2 |
| Total | 33 | 50 |

Table 3: Summary of Z declarations for Section 4.

5 Scheduler

section *Scheduler* parents *System*

This part of the specification is the scheduler.

Here, we declare the set of process IDs, the priority values, and the default number of credits a process receives when it is created.

Pid == \mathbb{N}
Priority == - 19 .. 19
Default_Credits == 10

The possible status that a process can hold.

Status ::= *pWaiting* | *pReady* | *pRunning*

A process consists of a process ID, a status, a number of credits, and a priority. Each process has a sequence of instructions to be executed on the assembler, with a pointer to the current instruction. The memory that a process can occupy is between a base and limit value. Instructions must only access memory with a value less than the limit, but they know nothing about the base value - this is added onto the memory index provided by the instruction when an instruction is executed. Each process also contains a stack and values for all registers, which are used to store values when the process is suspended.

Processes

$$\begin{array}{l}
pids : \mathbb{P} \text{ Pid} \\
status : \text{Pid} \leftrightarrow \text{Status} \\
credits : \text{Pid} \leftrightarrow \mathbb{N} \\
priority : \text{Pid} \leftrightarrow \text{Priority} \\
instructions : \text{Pid} \leftrightarrow (\text{seq } \text{Instruction}) \\
inst_pointer : \text{Pid} \leftrightarrow \mathbb{N}_1 \\
base, limit : \text{Pid} \leftrightarrow \text{WORD} \\
pregisters : \text{Pid} \leftrightarrow \text{REGISTERS} \\
pstack : \text{Pid} \leftrightarrow \text{Stack}[\text{WORD}] \\
\\
pids = \text{dom}(status) = \text{dom}(credits) = \text{dom}(priority) = \\
\quad \text{dom}(instructions) = \text{dom}(inst_pointer) = \text{dom}(base) = \\
\quad \text{dom}(limit) = \text{dom}(pstack) \\
\forall pid : pids \bullet inst_pointer(pid) \leq \#(instructions(pid)) \\
\forall pid : pids \bullet base(pid) + limit(pid) \leq mem_size
\end{array}$$

The *sort* function takes the credits and priorities of all processes, and returns a sequence of process IDS sorted firstly by their credits (the more credits a process has, the higher preference they get), and if the credits are equal, then their priority. If the priority is equal, then the order is non-deterministic.

$$\begin{array}{l}
sort : (\text{Pid} \leftrightarrow \mathbb{N}) \times (\text{Pid} \leftrightarrow \text{Priority}) \leftrightarrow \text{iseq } \text{Pid} \\
\\
sort = (\lambda credits : (\text{Pid} \leftrightarrow \mathbb{N}); priority : (\text{Pid} \leftrightarrow \text{Priority}) \mid \\
\quad \text{dom}(credits) = \text{dom}(priority) \bullet \\
\quad (\mu s : \text{iseq } \text{Pid} \mid \text{ran}(s) = \text{dom}(credits) \wedge \\
\quad \quad (\forall i : 1 \dots \# s - 1 \bullet \\
\quad \quad \quad credits(s(i)) > credits(s(i+1)) \vee \\
\quad \quad \quad (credits(s(i)) = credits(s(i+1))) \wedge \\
\quad \quad \quad priority(s(i)) > priority(s(i+1)))) \bullet s))
\end{array}$$

To interrupt a process during execution, the kernel must be in *kernel* mode.

Mode ::= user | kernel

For the scheduler, we track which mode the operating system is in, as well as declaring three “secondary” variables, *waiting*, *running*, and *ready*, to keep the sets of waiting running, and ready variables respectively. In fact, *ready* is a sequence, and is ordered based on the credits that each process has. A process with more credits will have a higher priority. This is fair scheduling, because at each timer interrupt (the *tick* operation specified below), the current process losses one credit, therefore, process spending a lot of time executing will eventually have a low priority.

| |
|-------------------------------------------------------------------------------------------------|
| <i>Scheduler</i> |
| <i>Processes</i> |
| <i>System</i> |
| <i>Stack</i> [<i>WORD</i>] |
| <i>mode</i> : <i>Mode</i> |
| <i>waiting, running</i> : $\mathbb{P} \text{ Pid}$ |
| <i>ready</i> : <i>iseq Pid</i> |
| $\# \text{ running} \leq 1$ |
| $\text{waiting} \cap \text{running} \cap \text{ranready} = \emptyset$ |
| $\text{waiting} \cup \text{running} \cup \text{ranready} = \text{pids}$ |
| $\text{waiting} = \{p : \text{pids} \mid (\text{status} \sim)(p\text{Waiting}) = p\}$ |
| $\text{running} = \{p : \text{pids} \mid (\text{status} \sim)(p\text{Running}) = p\}$ |
| $\text{ready} = \text{sort}((\text{waiting} \cup \text{running}) \triangleleft \text{credits},$ |
| $\quad (\text{waiting} \cup \text{running}) \triangleleft \text{priority})$ |
| $\forall r : \text{ran}(\text{ready}) \bullet \text{status}(r) = p\text{Ready}$ |
| $\forall r : \text{running} \bullet \text{credits}(r) > 0$ |

This uses semicolons as conjunctions for predicates, which conforms to the grammar in the ISO standard, but according to the list of differences between ZRM and ISO Z on Ian Toyn’s website, semicolons can no longer be used to conjoin predicates.

| |
|-------------------------------------------------------------------------------------------------|
| <i>InitScheduler</i> |
| <i>Scheduler</i> |
| <i>InitStack</i> [<i>WORD</i>] |
| <i>InitSystem</i> |
| $\text{pids} = \emptyset; \text{status} = \emptyset; \text{priority} = \emptyset$ |
| $\text{credits} = \emptyset; \text{instructions} = \emptyset; \text{inst_pointer} = \emptyset$ |
| $\text{waiting} = \emptyset; \text{running} = \emptyset; \text{ready} = \langle \rangle$ |
| $\text{base} = \{\}; \text{limit} = \{\}; \text{pregisters} = \{\}$ |
| $\text{mode} = \text{user}$ |

newProcess creates a new process with a unique process ID and a specified priority, and places this new process on the ready queue.

| | |
|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>create_new_process</i> | <hr/> $\Delta Scheduler$ $\Xi System$ $priority? : Priority$ $instructions? : seq Instruction$ $base?, limit? : WORD$ $pid! : Pid$ <hr/> $pid! \notin pids$ $status' = status \cup \{pid! \mapsto pReady\}$ $credits' = credits \cup \{pid! \mapsto DefaultCredits\}$ $priority' = priority \cup \{pid! \mapsto priority?\}$ $instructions' = instructions \cup \{pid! \mapsto instructions?\}$ $inst_pointer' = inst_pointer \cup \{pid! \mapsto 1\}$ $base' = base \cup \{pid! \mapsto base?\}$ $limit' = limit \cup \{pid! \mapsto limit?\}$ $pregisters' =$ $\quad pregisters \cup \{pid! \mapsto \{r : REGISTER \bullet$ $\quad \quad r \mapsto constant(min(WORD))\}\}$ $pstack' = pstack \cup \{pid! \mapsto (\langle stack == \rangle)\}$ $pids' = pids \cup \{pid!\}$ <hr/> |
|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

We define a schema that contains only the variables that do not change when a reschedule occurs.

$$RescheduleChange == \\ Scheduler \setminus (status, running, ready, waiting, credits)$$

A reschedule occurs when all ready processes have no credits. Every process, not just the ready processes, have their credits re-calculated using the formula $credits = credits/2 + priority$. This implies that the ready process with the highest priority will be the next process executed.

| | |
|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>reschedule</i> | <hr/> $\Delta Scheduler$ $\Xi RescheduleChange$ <hr/> $ready \neq \emptyset$ $\forall r : ran(ready) \bullet credits(r) = 0 \Rightarrow$ $\quad credits' = \{p : pids \bullet p \mapsto (credits(p) \div 2) + priority(p)\} \wedge$ $\quad status' = status$ $\neg (\forall r : ran(ready) \bullet credits(r) = 0) \Rightarrow$ $\quad status' = status \oplus \{head(ready) \mapsto pRunning\} \wedge$ $\quad credits' = credits$ <hr/> |
|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

We declare a new schema that contains only the state variables that do not change when a status change occurs.

$$\text{StatusChange} == \text{Scheduler} \setminus (\text{status}, \text{running}, \text{waiting}, \text{ready}, \text{registers}, \text{pregisters}, \text{pstack})$$

Interrupts the currently executing process if the new process is of a higher priority then the current process and the kernel is in *kernel* mode.

$$\begin{array}{l} \text{interrupt} \text{ ---} \\ \Delta \text{Scheduler} \\ \Xi \text{StatusChange} \\ \text{create_new_process} \\ \hline \text{mode} = \text{kernel} \\ \text{running} = \emptyset \vee (\exists p : \text{running} \bullet \text{priority?} \geq \text{priority}(p)) \\ \exists r : \text{running} \bullet \\ \quad \text{status}' = \text{status} \oplus \{ \text{pid!} \mapsto p\text{Running}, r \mapsto p\text{Ready} \} \wedge \\ \quad \text{pregisters}' = \text{pregisters} \oplus \{ r \mapsto \text{registers} \} \wedge \\ \quad \theta \text{Stack}' = \text{pstack}(r) \\ \quad \text{registers}' = \text{pregisters}(\text{pid!}) \end{array}$$

Remove the currently running process and put it back in the ready queue.

$$\begin{array}{l} \text{remove_running_process} \text{ ---} \\ \Delta \text{Scheduler} \\ \Xi \text{StatusChange} \\ \hline \exists \text{pid} == (\mu r : \text{running}) \bullet \\ \quad \text{status}' = \text{status} \oplus \{ \text{pid} \mapsto p\text{Ready} \} \wedge \\ \quad \text{pregisters}' = \text{pregisters} \oplus \{ \text{pid} \mapsto \text{registers} \} \wedge \\ \quad \text{pstack}' = \text{pstack} \oplus \{ \text{pid} \mapsto \theta \text{Stack}' \} \end{array}$$

A process becomes blocked if it is waiting on a resource such as an IO device, or waiting on another process

$$\text{block_process} == \text{remove_running_process} \circ \text{reschedule}$$

We declare a schema containing only the variables that change for an unblock.

$$\text{UnblockProcessChange} == \text{Scheduler} \setminus (\text{status}, \text{running}, \text{ready}, \text{waiting})$$

Unblocks a process that is blocked by another process.

| |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\begin{array}{l} \text{unblock_process} \\ \hline \Delta \text{Scheduler} \\ \Xi \text{UnblockProcessChange} \\ pid? : \text{Pid} \end{array}$ |
| $\begin{array}{l} pid? \in pids \\ status(pid?) = p \text{Waiting} \\ running = \emptyset \Leftrightarrow status' = status \oplus \{pid? \mapsto p \text{Running}\} \\ running \neq \emptyset \Leftrightarrow status' = status \oplus \{pid? \mapsto p \text{Ready}\} \end{array}$ |

Remove a process from the system

| |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\begin{array}{l} \text{remove_process} \\ \hline \Delta \text{Scheduler} \\ \Xi \text{Stack}[\text{WORD}] \\ \Xi \text{System} \\ pid? : \text{Pid} \end{array}$ |
| $\begin{array}{l} pid? \in pids \\ pids' = pids \setminus \{pid?\} \\ status' = \{pid?\} \triangleleft status \\ credits' = \{pid?\} \triangleleft credits \\ priority' = \{pid?\} \triangleleft priority \\ instructions' = \{pid?\} \triangleleft instructions \\ inst_pointer' = \{pid?\} \triangleleft inst_pointer \\ base' = \{pid?\} \triangleleft base \\ limit' = \{pid?\} \triangleleft limit \\ pregisters' = \{pid?\} \triangleleft pregisters \\ pstack' = \{pid?\} \triangleleft pstack \end{array}$ |

Update the details in the process table when each instruction is executed, as well as communicate the current instruction and the base value for the current process.

$$\text{ChangeInstPointer} == \text{Scheduler} \setminus (\text{inst_pointer})$$

| |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>update_process_table</i> |
| $\Delta \text{Scheduler}$ $inst! : \text{Instruction}$ $base! : \text{WORD}$ |
| $running \neq \emptyset$ $(\exists pid == (\mu r : running) \bullet$ $\quad inst! = head(instructions(pid)) \wedge$ $\quad base! = base(pid) \wedge$ $\quad (inst_pointer(pid) = \#(instructions(pid)) \Rightarrow$ $\quad \quad remove_process[pid/pid?]) \wedge$ $\quad inst_pointer(pid) < \#(instructions(pid)) \Rightarrow$ $\quad \quad inst_pointer' =$ $\quad \quad \quad inst_pointer \oplus \{pid \mapsto inst_pointer(pid) + 1\})$ $\quad \theta \text{ ChangeInstPointer} = \theta \text{ ChangeInstPointer}'$ |

$next == exec_inst \gg update_process_table \circ$
 $([\Delta \text{Scheduler} \mid running = \emptyset] \wedge reschedule) \vee$
 $([\Xi \text{Scheduler} \mid running \neq \emptyset])$

$idle0 == \neg \text{pre } next$

| |
|---------------------------------------------------------------------------------|
| <i>idle</i> |
| $\Xi \text{Scheduler}$ $inst? : \text{Instruction}$ $base? : \text{WORD}$ |
| <i>idle0</i> |

$tick == next \vee idle$

$\vdash? (\forall n : \mathbb{N}_1 \bullet n > 0)$

$[X] \vdash? \forall x : \mathbb{P} X \bullet \# x \leq 1 \Leftrightarrow singleton\ x$

theorem PreconditionCheck
 $\forall \text{Scheduler} \bullet \text{pre } update_process_table$

| Declarations | This section | Globally |
|-------------------------|---------------------|-----------------|
| Unboxed items | 21 | 52 |
| Axiomatic definitions | 1 | 3 |
| Generic axiomatic defs. | 0 | 1 |
| Schemas | 11 | 25 |
| Generic schemas | 0 | 2 |
| Total | 33 | 83 |

Table 4: Summary of Z declarations for Section 5.