

## 1 Introduction

This is a specification of a simple scheduler and assembler. The system contains a set of registers and a block of memory. Processes can be created, with each containing a sequence of instructions that are executed on the system. The instruction format is a simplified format of the Intel x86 architecture. Processes are scheduled based on the credit system that is found in the Linux 2.0 kernel.

## 2 Stack

This specification was written as a test spec for the CZT project. As a result, there are parts that may appear to be specified in a strange way - this is to test out the tools on a large set of Z.

**section** *Stack* **parents** *standard\_toolkit*

A generic stack.

$Stack[X] == [stack : seq X]$   
 $InitStack[X] == [Stack[X] \mid stack = \emptyset]$

$PushStack[X]$ _____ $\Delta Stack[X]$ $x? : X$
$stack' = stack \frown \langle x? \rangle$

$PopStack[X]$ _____ $\Delta Stack[X]$ $x! : X$
$stack' \frown \langle x! \rangle = stack$

Ok, lets see the value of 3 unboxed items in Section 2!

## 3 Definitions

**section** *Definitions* **parents** *standard\_toolkit*

Declarations	This Section	Globally
Unboxed items	3	3
Axiomatic definitions	0	0
Generic axiomatic defs.	0	0
Schemas	0	0
Generic schemas	2	2
<b>Total</b>	<b>5</b>	<b>5</b>

Table 1: Summary of Z declarations for Section 2.

Firstly, we define some basic types and functions that are used throughout the specification.

*singleton* is the set of all sets whose size is less than or equal to 1. This is included only to have a generic axiom definition.

**relation**(*singleton* \_)

[X]	
<i>singleton</i> _ : $\mathbb{P}(\mathbb{P} X)$	
$\forall s : \mathbb{P} X \bullet \text{singleton } s \Leftrightarrow \# s \leq 1$	

The basic type of this system is a word, which specifically, is an unsigned octet. An unsigned word is used so references to memory etc a 1-relative.

**WORD** == 0 .. 255

Then, we define the size of the memory block, and give it a value for animation purposes.

<i>mem_size</i> : <b>WORD</b>
<i>mem_size</i> = 100

A *LABEL* is used to label instructions for *jump* instructions etc, although 'jump' hasn't been specified yet.

[*LABEL*]

Now we define the different instructions, as well as their operands. A *CONSTANT* is used both as a constant value, as well as a memory reference for load and store instructions.

$INST\_NAME ::=$   
 $add \mid sub \mid divide \mid mult \mid push \mid pop \mid load \mid store \mid loadConst \mid print$   
 $OPERAND ::= AX \mid BX \mid CX \mid DX \mid constant\langle\langle WORD \rangle\rangle$   
 $REGISTER == \{AX, BX, CX, DX\}$   
 $CONSTANT == OPERAND \setminus REGISTER$

An instruction is specified as a instruction name, a sequence of operands, and optionally, a label.

$Instruction$ $label : \mathbb{P} LABEL$ $name : INST\_NAME$ $params : seq OPERAND$ $singleton label$
---

Declarations	This Section	Globally
Unboxed items	9	12
Axiomatic definitions	1	1
Generic axiomatic defs.	1	1
Schemas	1	1
Generic schemas	0	2
<b>Total</b>	<b>12</b>	<b>17</b>

Table 2: Summary of Z declarations for Section 3.

## 4 System

**section** *System* **parents** *Definitions, Stack*

The system consists of a set of registers, and a block of memory. There is also a buffer for displaying output.

$REGISTERS == REGISTER \rightarrow OPERAND$   
 $MEMORY == 1 .. mem\_size \rightarrow WORD$

---

*System*

*registers* : *REGISTERS*  
*memory* : *MEMORY*  
*output* : seq *WORD*

---

Initially, all registers and memory hold the minimum *WORD* value. The output buffer is empty.

---

*InitSystem*

---

*System*

---

*registers* = {*r* : *REGISTER* • *r* ↦ *constant*(*min*(*WORD*))}  
*memory* = {*m* : 1 .. *mem\_size* • *m* ↦ *min*(*WORD*)}  
*output* = ⟨⟩

---

The system can have arithmetic and memory instructions.

*Arith\_Inst* == [*Instruction* | #*params* = 2 ∧ *params*(1) ∈ *REGISTER*]  
*Add\_Inst* == [*Arith\_Inst* | *name* = *add*]  
*Sub\_Inst* == [*Arith\_Inst* | *name* = *sub*]  
*Mult\_Inst* == [*Arith\_Inst* | *name* = *mult*]  
*Div\_Inst* == [*Arith\_Inst* | *name* = *divide*]

*Memory\_Inst* == [*Instruction* | #*params* = 2 ∧ *params*(1) ∈ *REGISTER*  
 ∧ *params*(2) ∈ *CONSTANT*]  
*Load\_Inst* == [*Memory\_Inst* | *name* = *load*]  
*LoadConst\_Inst* == [*Memory\_Inst* | *name* = *loadConst*]  
*Store\_Inst* == [*Memory\_Inst* | *name* = *store*]

A print instruction prints the value of a register.

*Print\_Inst* == [*Instruction* | #*params* = 1]

*val* maps constants to their value, and *dereference* dereferences the value of a register, transitively if required.

---

*val* : *CONSTANT* → *WORD*  
*dereference* : *OPERAND* × *REGISTERS* → *WORD*

---

∀ *c* : *CONSTANT* •  
 (∃ *n* : *WORD* • *c* = *constant*(*n*) ∧ *val*(*c*) = *n*)  
 ∀ *a* : *OPERAND*; *r* : *REGISTERS* •  
*dereference*(*a*, *r*) =  
 if *a* ∈ *REGISTER* then *dereference*(*r*(*a*), *r*) else *val*(*a*)

---

The specification of the arithmetic instructions.

<i>Add</i>
$\Delta_{System}$ <i>Add_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 + o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Sub</i>
$\Delta_{System}$ <i>Sub_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 - o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Mult</i>
$\Delta_{System}$ <i>Mult_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 * o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

<i>Div</i>
$\Delta_{System}$ <i>Div_Inst</i>
$\begin{aligned} &\exists o_1 == \text{dereference}(\text{params}(1), \text{registers}); \\ &\quad o_2 == \text{dereference}(\text{params}(2), \text{registers}) \bullet \\ &\quad \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_1 \text{ div } o_2)\} \\ &\text{memory}' = \text{memory} \\ &\text{output}' = \text{output} \end{aligned}$

The **load** operation loads a constant from memory. The second parameter is an index to the memory location from which the constant is loaded.

<i>Load</i>
$\Delta System$
<i>Load_Inst</i>
$\begin{aligned} \exists o_2 == & \text{val}(\text{params}(2)) \bullet \\ & \text{registers}' = \text{registers} \oplus \\ & \quad \{\text{params}(1) \mapsto \text{constant}(\text{memory}(o_2))\} \\ \text{memory}' = & \text{memory} \\ \text{output}' = & \text{output} \end{aligned}$

**loadConst** loads a constant into a register. The second parameter the constant to be loaded.

<i>Load_Const</i>
$\Delta System$
<i>LoadConst_Inst</i>
$\begin{aligned} \exists o_2 == & \text{val}(\text{params}(2)) \bullet \\ & \text{registers}' = \text{registers} \oplus \{\text{params}(1) \mapsto \text{constant}(o_2)\} \\ \text{memory}' = & \text{memory} \\ \text{output}' = & \text{output} \end{aligned}$

Store the value of a register in memory.

<i>Store</i>
$\Delta System$
<i>Store_Inst</i>
$\begin{aligned} \exists o_1 == & \text{dereference}(\text{params}(1), \text{registers}); \\ & o_2 == \text{val}(\text{params}(2)) \bullet \\ & \quad \text{memory}' = \text{memory} \oplus \{o_2 \mapsto o_1\} \\ \text{registers}' = & \text{registers} \\ \text{output}' = & \text{output} \end{aligned}$

<i>Print</i>
$\Xi System$
<i>Print_Inst</i>
$\begin{aligned} \text{output}' = & \text{output} \frown \langle \text{dereference}(\text{params}(1), \text{registers}) \rangle \\ \text{registers}' = & \text{registers} \\ \text{memory}' = & \text{memory} \end{aligned}$

$$\begin{aligned}
Stack\_Inst &== [Instruction \mid \#params = 1] \\
Push\_Inst &== [Stack\_Inst \mid name = push] \\
Pop\_Inst &== [Stack\_Inst \mid name = pop]
\end{aligned}$$

The specification of the stack instructions on the system.

$ \begin{aligned} &Push0 \\ &\Xi System \\ &PushStack[WORD] \\ &Push\_Inst \end{aligned} $
$x? = dereference(params(1), registers)$

$ \begin{aligned} &Pop0 \\ &\Delta System \\ &PopStack[WORD] \\ &Pop\_Inst \end{aligned} $
$ \begin{aligned} registers' &= registers \oplus \{params(1) \mapsto constant(x!)\} \\ memory' &= memory \\ output' &= output \end{aligned} $

$$\begin{aligned}
Push &== Push0 \upharpoonright [System; Stack[WORD]] \\
Pop &== Pop0 \upharpoonright [System; Stack[WORD]]
\end{aligned}$$

This executes an instruction on the on the system. *inst?* is the instruction to execute, and *base?* is the base memory value of the executing process. If the instruction is a **load** or **store** instruction, the memory reference must offset using the base value.

<i>exec_inst</i>
$\Delta System$
<i>inst?</i> : <i>Instruction</i>
<i>base?</i> : 1 .. <i>mem_size</i>
$\exists label : \mathbb{P} LABEL; name : INST\_NAME; params : seq OPERAND \mid$ $label = inst?.label \wedge name = inst?.name \wedge$ $params = inst?.params \bullet$ $Add \vee Sub \vee Mult \vee Div \vee$ $Print \vee Load\_Const \vee$ $name \in \{load, store\} \Rightarrow (\exists p : seq OPERAND \mid$ $p = \langle params(1),$ $constant(val(params(2)) + base?) \rangle \bullet$ $Load[p/params] \vee Store[p/params])$

Declarations	This Section	Globally
Unboxed items	19	31
Axiomatic definitions	1	2
Generic axiomatic defs.	0	1
Schemas	13	14
Generic schemas	0	2
<b>Total</b>	<b>33</b>	<b>50</b>

Table 3: Summary of Z declarations for Section 4.

## 5 Scheduler

### section *Scheduler* parents *System*

This part of the specification is the scheduler.

Here, we declare the set of process IDs, the priority values, and the default number of credits a process receives when it is created.

*Pid* ==  $\mathbb{N}$   
*Priority* == -19 .. 19  
*Default\_Credits* == 10

The possible status that a process can hold.

*Status* ::= *pWaiting* | *pReady* | *pRunning*



A process consists of a process ID, a status, a number of credits, and a priority. Each process has a sequence of instructions to be executed on the assembler, with a pointer to the current instruction. The memory that a process can occupy is between a base and limit value. Instructions must only access memory with a value less than the limit, but they know nothing about the base value - this is added onto the memory index provided by the instruction when an instruction is executed. Each process also contains a stack and values for all registers, which are used to store values when the process is suspended.

#### *Processes*

$$\begin{array}{l}
pids : \mathbb{P} \text{ Pid} \\
status : \text{Pid} \leftrightarrow \text{Status} \\
credits : \text{Pid} \leftrightarrow \mathbb{N} \\
priority : \text{Pid} \leftrightarrow \text{Priority} \\
instructions : \text{Pid} \leftrightarrow (\text{seq } \text{Instruction}) \\
inst\_pointer : \text{Pid} \leftrightarrow \mathbb{N}_1 \\
base, limit : \text{Pid} \leftrightarrow \text{WORD} \\
pregisters : \text{Pid} \leftrightarrow \text{REGISTERS} \\
pstack : \text{Pid} \leftrightarrow \text{Stack}[\text{WORD}] \\
\hline
pids = \text{dom}(status) = \text{dom}(credits) = \text{dom}(priority) = \\
\quad \text{dom}(instructions) = \text{dom}(inst\_pointer) = \text{dom}(base) = \\
\quad \text{dom}(limit) = \text{dom}(pstack) \\
\forall pid : pids \bullet inst\_pointer(pid) \leq \#(instructions(pid)) \\
\forall pid : pids \bullet base(pid) + limit(pid) \leq mem\_size
\end{array}$$

The *sort* function takes the credits and priorities of all processes, and returns a sequence of process IDS sorted firstly by their credits (the more credits a process has, the higher preference they get), and if the credits are equal, then their priority. If the priority is equal, then the order is non-deterministic.

$$\begin{array}{l}
sort : (\text{Pid} \leftrightarrow \mathbb{N}) \times (\text{Pid} \leftrightarrow \text{Priority}) \leftrightarrow \text{iseq } \text{Pid} \\
\hline
sort = (\lambda credits : (\text{Pid} \leftrightarrow \mathbb{N}); priority : (\text{Pid} \leftrightarrow \text{Priority}) \mid \\
\quad \text{dom}(credits) = \text{dom}(priority) \bullet \\
\quad (\mu s : \text{iseq } \text{Pid} \mid \text{ran}(s) = \text{dom}(credits) \wedge \\
\quad \quad (\forall i : 1 \dots \# s - 1 \bullet \\
\quad \quad \quad credits(s(i)) > credits(s(i+1)) \vee \\
\quad \quad \quad (credits(s(i)) = credits(s(i+1))) \wedge \\
\quad \quad \quad priority(s(i)) > priority(s(i+1)))) \bullet s))
\end{array}$$

To interrupt a process during execution, the kernel must be in *kernel* mode.

*Mode ::= user | kernel*

For the scheduler, we track which mode the operating system is in, as well as declaring three “secondary” variables, *waiting*, *running*, and *ready*, to keep the sets of waiting running, and ready variables respectively. In fact, *ready* is a sequence, and is ordered based on the credits that each process has. A process with more credits will have a higher priority. This is fair scheduling, because at each timer interrupt (the *tick* operation specified below), the current process losses one credit, therefore, process spending a lot of time executing will eventually have a low priority.

#### *Scheduler*

*Processes*

*System*

*Stack*[WORD]

*mode* : Mode

*waiting, running* :  $\mathbb{P} \text{Pid}$

*ready* : *iseq* *Pid*

$\# \text{running} \leq 1$

$\text{waiting} \cap \text{running} \cap \text{ran ready} = \emptyset$

$\text{waiting} \cup \text{running} \cup \text{ran ready} = \text{pids}$

$\text{waiting} = \{p : \text{pids} \mid (\text{status} \sim)(p\text{Waiting}) = p\}$

$\text{running} = \{p : \text{pids} \mid (\text{status} \sim)(p\text{Running}) = p\}$

$\text{ready} = \text{sort}((\text{waiting} \cup \text{running}) \triangleleft \text{credits},$

$(\text{waiting} \cup \text{running}) \triangleleft \text{priority})$

$\forall r : \text{ran}(\text{ready}) \bullet \text{status}(r) = p\text{Ready}$

$\forall r : \text{running} \bullet \text{credits}(r) > 0$

This uses semicolons as conjunctions for predicates, which conforms to the grammar in the ISO standard, but according to the list of differences between ZRM and ISO Z on Ian Toyn’s website, semicolons can no longer be used to conjoin predicates.

#### *InitScheduler*

*Scheduler*

*InitStack*[WORD]

*InitSystem*

$\text{pids} = \emptyset ; \text{status} = \emptyset ; \text{priority} = \emptyset$

$\text{credits} = \emptyset ; \text{instructions} = \emptyset ; \text{inst\_pointer} = \emptyset$

$\text{waiting} = \emptyset ; \text{running} = \emptyset ; \text{ready} = \langle \rangle$

$\text{base} = \{\} ; \text{limit} = \{\} ; \text{pregisters} = \{\}$

$\text{mode} = \text{user}$

*newProcess* creates a new process with a unique process ID and a specified priority, and places this new process on the ready queue.

*create\_new\_process*

$\Delta$ *Scheduler*

$\Xi$ *System*

*priority?* : *Priority*

*instructions?* : seq *Instruction*

*base?*, *limit?* : *WORD*

*pid!* : *Pid*

$pid! \notin pids$

$status' = status \cup \{pid! \mapsto pReady\}$

$credits' = credits \cup \{pid! \mapsto DefaultCredits\}$

$priority' = priority \cup \{pid! \mapsto priority?\}$

$instructions' = instructions \cup \{pid! \mapsto instructions?\}$

$inst\_pointer' = inst\_pointer \cup \{pid! \mapsto 1\}$

$base' = base \cup \{pid! \mapsto base?\}$

$limit' = limit \cup \{pid! \mapsto limit?\}$

$pregisters' =$

$\quad pregisters \cup \{pid! \mapsto \{r : REGISTER \bullet$   
 $\quad \quad r \mapsto constant(min(WORD))\}\}$

$pstack' = pstack \cup \{pid! \mapsto (\langle stack == \rangle)\}$

$pids' = pids \cup \{pid!\}$

We define a schema that contains only the variables that do not change when a reschedule occurs.

*RescheduleChange* ==

*Scheduler*  $\setminus$  (*status*, *running*, *ready*, *waiting*, *credits*)

A reschedule occurs when all ready processes have no credits. Every process, not just the ready processes, have their credits re-calculated using the formula  $credits = credits/2 + priority$ . This implies that the ready process with the highest priority will be the next process executed.

*reschedule*

$\Delta$ *Scheduler*

$\Xi$ *RescheduleChange*

$ready \neq \emptyset$

$\forall r : \text{ran}(ready) \bullet credits(r) = 0 \Rightarrow$

$\quad credits' = \{p : pids \bullet p \mapsto (credits(p) \text{ div } 2) + priority(p)\} \wedge$   
 $\quad status' = status$

$\neg (\forall r : \text{ran}(ready) \bullet credits(r) = 0) \Rightarrow$

$\quad status' = status \oplus \{head(ready) \mapsto pRunning\} \wedge$   
 $\quad credits' = credits$

We declare a new schema that contains only the state variables that do not change when a status change occurs.

$$\text{StatusChange} == \text{Scheduler} \setminus (\text{status}, \text{running}, \text{waiting}, \text{ready}, \text{registers}, \text{pregisters}, \text{pstack})$$

Interrupts the currently executing process if the new process is of a higher priority then the current process and the kernel is in *kernel* mode.

$$\begin{array}{l} \text{interrupt} \\ \hline \Delta \text{Scheduler} \\ \exists \text{StatusChange} \\ \text{create\_new\_process} \\ \hline \text{mode} = \text{kernel} \\ \text{running} = \emptyset \vee (\exists p : \text{running} \bullet \text{priority?} \geq \text{priority}(p)) \\ \exists r : \text{running} \bullet \\ \quad \text{status}' = \text{status} \oplus \{pid! \mapsto p\text{Running}, r \mapsto p\text{Ready}\} \wedge \\ \quad \text{pregisters}' = \text{pregisters} \oplus \{r \mapsto \text{registers}\} \wedge \\ \quad \theta \text{Stack}' = \text{pstack}(r) \\ \quad \text{registers}' = \text{pregisters}(pid!) \end{array}$$

Remove the currently running process and put it back in the ready queue.

$$\begin{array}{l} \text{remove\_running\_process} \\ \hline \Delta \text{Scheduler} \\ \exists \text{StatusChange} \\ \hline \exists pid == (\mu r : \text{running}) \bullet \\ \quad \text{status}' = \text{status} \oplus \{pid \mapsto p\text{Ready}\} \wedge \\ \quad \text{pregisters}' = \text{pregisters} \oplus \{pid \mapsto \text{registers}\} \wedge \\ \quad \text{pstack}' = \text{pstack} \oplus \{pid \mapsto \theta \text{Stack}'\} \end{array}$$

A process becomes blocked if it is waiting on a resource such as an IO device, or waiting on another process

$$\text{block\_process} == \text{remove\_running\_process} \circ \text{reschedule}$$

We declare a schema containing only the variables that change for an unblock.

$$\text{UnblockProcessChange} == \text{Scheduler} \setminus (\text{status}, \text{running}, \text{ready}, \text{waiting})$$

Unblocks a process that is blocked by another process.

<i>unblock_process</i>	_____
$\Delta\text{Scheduler}$	
$\Xi\text{UnblockProcessChange}$	
$pid? : \text{Pid}$	
$pid? \in pids$	
$status(pid?) = pWaiting$	
$running = \emptyset \Leftrightarrow status' = status \oplus \{pid? \mapsto pRunning\}$	
$running \neq \emptyset \Leftrightarrow status' = status \oplus \{pid? \mapsto pReady\}$	

Remove a process from the system

<i>remove_process</i>	_____
$\Delta\text{Scheduler}$	
$\Xi\text{Stack}[\text{WORD}]$	
$\Xi\text{System}$	
$pid? : \text{Pid}$	
$pid? \in pids$	
$pids' = pids \setminus \{pid?\}$	
$status' = \{pid?\} \triangleleft status$	
$credits' = \{pid?\} \triangleleft credits$	
$priority' = \{pid?\} \triangleleft priority$	
$instructions' = \{pid?\} \triangleleft instructions$	
$inst\_pointer' = \{pid?\} \triangleleft inst\_pointer$	
$base' = \{pid?\} \triangleleft base$	
$limit' = \{pid?\} \triangleleft limit$	
$pregisters' = \{pid?\} \triangleleft pregisters$	
$pstack' = \{pid?\} \triangleleft pstack$	

Update the details in the process table when each instruction is executed, as well as communicate the current instruction and the base value for the current process.

$$\text{ChangeInstPointer} == \text{Scheduler} \setminus (inst\_pointer)$$

$update\_process\_table$ $\Delta Scheduler$ $inst! : Instruction$ $base! : WORD$
$running \neq \emptyset$ $(\exists pid == (\mu r : running) \bullet$ $\quad inst! = head(instructions(pid)) \wedge$ $\quad base! = base(pid) \wedge$ $\quad (inst\_pointer(pid) = \#(instructions(pid)) \Rightarrow$ $\quad \quad remove\_process[pid/pid?]) \wedge$ $\quad inst\_pointer(pid) < \#(instructions(pid)) \Rightarrow$ $\quad \quad inst\_pointer' =$ $\quad \quad \quad inst\_pointer \oplus \{pid \mapsto inst\_pointer(pid) + 1\})$ $\theta ChangeInstPointer = \theta ChangeInstPointer')$

$next == exec\_inst \gg update\_process\_table \circ$   
 $([\Delta Scheduler \mid running = \emptyset] \wedge reschedule) \vee$   
 $([\exists Scheduler \mid running \neq \emptyset])$

$idle0 == \neg \mathbf{pre} next$

$idle$ $\exists Scheduler$ $inst? : Instruction$ $base? : WORD$
$idle0$

$tick == next \vee idle$

$\vdash? (\forall n : \mathbb{N}_1 \bullet n > 0)$

$[X] \vdash? \forall x : \mathbb{P} X \bullet \#x \leq 1 \Leftrightarrow singleton\ x$

**theorem** PreconditionCheck  
 $\forall Scheduler \bullet \mathbf{pre} update\_process\_table$

<b>Declarations</b>	<b>This Section</b>	<b>Globally</b>
Unboxed items	21	52
Axiomatic definitions	1	3
Generic axiomatic defs.	0	1
Schemas	11	25
Generic schemas	0	2
<b>Total</b>	<b>33</b>	<b>83</b>

Table 4: Summary of Z declarations for Section 5.